



# UPPSALA UNIVERSITET

## Parallel Quicksort Implementation with OpenMP

Yuanjing Yang

July 11, 2025

## 1 Introduction

Quicksort is one of the fastest sorting algorithms, which is division and conquer-based sorting algorithm that picks an element as a pivot and partitions the given array around the picked pivot, the lower values are on the left of the pivot and higher values on the right of it, and then recursively sort it until the sub\_array are too small to be sorted. The worst case scenario for quicksort that the pivot element is either the highest or lowest value in every sub\_array is  $O(n^2)$ , and on average, the time complexity is  $O(N \log N)$ , the recursive part of quicksort is the reason [1], the good pick of the pivot element, the array will be split in half somewhat evenly each time the algorithm calls itself.

The divide\_and\_conquer nature of Quicksort makes it a natural candidate for parallelization. Parallel computing offers an effective way to reduce execution time and improve performance. By breaking down the sorting problem into smaller, independent sub\_problems that can be processed concurrently.

## 2 Problem description

The primary objective is to sort a large array of integers using parallel computing techniques. The program accept three input arguments as follows: **project n num.threads**

where the input arguments have the following meaning:

**n** is the size of array

**n\_threads** is the number of threads to use, and it must be a power of two and greater than zero.

In parallel quicksort, it involves selecting a pivot, partitioning data around pivot, and then recursively sorting the sub-arrays, and in parallel context, these steps are distributed among multiple threads to achieve concurrent execution and accelerate the sorting process [2].

### 2.1 Solution method

The implemented parallel quicksort is optimized by shared-memory systems using Openmp. The main idea is to distribute the sorting tasks among available threads, pivot-based partitioning, and data exchanging.

Algorithm Overview:

1. Selecting a Global Pivot:

In **select\_pivot** function, each thread sort its array and computes the local median, and collect it to **global\_medians[group][locid]**, where **group** identifies the current logical thread group(**myid/current\_group\_size0**, and **locid** is the thread's ID within that group(**myid % current\_group\_size0**). And only the thread with **locid==0** collects all local medians within its group, sorts them and selects the median of these medians as the global pivot for the current recursion level. The chosen pivot is sorted in **global\_pivots[group]** [3].

(2) Local Data partitioning:

In **findsplit** function, it return the **split** index. In this part, for avoiding race conditions and improve temporal locality, each thread has its private copy of the buffer, the elements are written to temporary buffer based on pivot condition, and then copy back to the original array.

(3) Parallel Data exchange:

In **exchange\_data** function, I was inspired by the MPI communication pattern by exchanging boundary information between partitions. For lower half of current group, it keeps the elements less than or equal to the pivot, and sends the elements greater than the pivot, and the upper half is opposite, it keeps the elements greater than the pivot and sends the elements less than or equal to pivot, and then **send\_scr** data is copied into **thread\_tem\_buffers[myid]**, it similar to the concept of message buffers in MPI. Each thread identified its partner thread for exchanging data. The partner thread repeat the operation above. Here I used double-buffering system [4], which can manage data efficiently and reduce memory allocation overhead during exchange, and the **buffer\_track[myid]** array keeps track of the currently active buffer and the **target\_buffer\_index** determines the inactive buffer that merged data will be written, the **target\_buffer** pointer is then set to **thread\_buffers\_A[myid]** or **thread\_buffers\_B[myid]**. [5]

(4) Recursive calls:

The **global\_sort** function is called recursively, and the **current\_gourp\_size** is halved in each recursive call, and eventually leading to **current\_group\_size** **:=1**, and each thread performed a final **sort\_array** [6].

## 2.2 Experiment

1. Correctness verification:

In **main** function, I used **is\_sorted** function to confirm that all elements are in ascending order, the example output:

```
yuya8334@yyj-172525:/mnt/d/project$ ./project 1000000 4
OK! Array is sorted.
Time: 0.0327s
```

Figure 1: Example output(N=1,000,000)

2. Evaluation of Performance:

As the table showed below, we can see the execution time increases with the array sized for a fixed number of threads, this indicates that the complexity  $O(N \log N)$  and this algorithm scales well with increasing input array size.

To evaluate the performance of parallel efficiency, I tested the situation that fixed array size with different number of threads.

$$\text{Speedup} = \frac{T_1}{T_n}$$

Table 1: Execution Time for  $P = 4$  Threads

Array Size ( $N$ )	Time (s)
1,000,000	0.0344
5,000,000	0.2039
10,000,000	0.3985
50,000,000	2.1566
100,000,000	4.4815

If the speedup is ideal, then with  $n$  threads, the execution time should be

$$T_n = \frac{T_1}{n}$$

From the figure below, we can see that while speedup reached 5.29x at 16 threads, the efficiency dropped to 0.33, indicating that each additional thread contributed less effectively to overall performance. Within the increase of threads, the speedup is achieved, but it is not ideal, and the efficiency is decreasing.

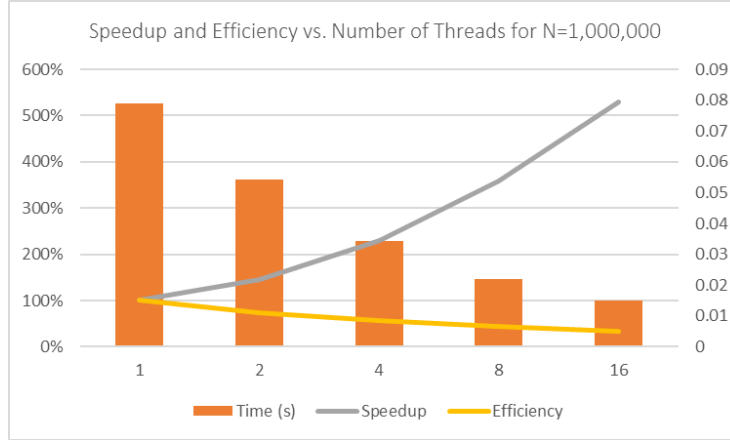


Figure 2: Speedup and Efficiency vs. Number of Threads(N=1,000,000)

There has two reasons for this results, one is communication overhead, even this algorithm used double-buffering system, but the data exchanging between threads possibly leads to cache coherence overhead. One is Synchronization Overhead, even the pivot strategy creates balanced partitions, but with the use of synchronization(**pragma omp barrier**), all threads need to wait until the slowest thread reaches that points.

### 3 Conclusion

This project implemented and evaluated an OpenMP-based parallel Quicksort algorithm, which is a divided-and-conquer strategy with double-buffering system

and data exchanging, parallelizing the sorting processes on shared-memory. The result shows the algorithm's correctness, and it achieves a speedup, reducing the execution time. However, the speedup didn't fit the expectation perfectly, it can be considered about influence of communication and Synchronization overheads. Here are some ideas for improvement:

1. Using MPI in **exchange\_data** function.
2. Using Dynamic scheduling [7] for better load balancing(I have tried this before, but it hard to match with the strategy of data exchanging, and it causes more complex race conditions, leading to a bad performance).

## 4 Declaration of AI

In my project, I used Chatgpt and deepseek to write some repeated and basic codes, debugging and make a improvement of my code. I write the original code frame and tell AI my the ideas, for example, I studied parallel and distribution programming in this period, and learn the use of MPI [6], so I plan to use it in this project, chatgpt help me Outlines the process of converting from MPI to OpenMP(in exchanging data part). AI help me to find some related materials, I found a way to improve performance(double buffer system) and looked for the bibliographic on online library, it increases the efficiency a lot, I don't need to waste a lot time in searching method. AI also helped me to check the debug information, especially when I used gdb and valgrind, it helps me find the errors quickly, and AI also help me to modify my code based on my results, I always get some useful feedback from AI.

## References

- [1] W3Schools. Dsa - quick sort algorithm, 2025. Accessed: July 11, 2025.
- [2] GeeksforGeeks. Implementation of quick sort using mpi, omp and posix thread, 2025. Accessed: July 11, 2025.
- [3] year = 2025 note = Accessed: 4 11, 2025 url = Quicksort.pdf Jarmo Rantakokko, title = Parallel Quick Sort.
- [4] Reddit User Community. How does a double buffer allow things to be drawn smoothly?, 2022. Accessed: July 11, 2025.
- [5] Eko Dwi Nugroho, Ilham Firman Ashari, Muhammad Nashrullah, Muhammad Habib Algifari, and Miranti Verdiana. Comparative analysis of openmp and mpi parallel computing implementations in team sort algorithm. *Journal of Applied Informatics and Computing*, 7(2):141–149, 2023.
- [6] Peter S. Pacheco. *An introduction to parallel programming*. Elsevier/Morgan Kaufmann, Amsterdam ;, 2011.
- [7] Kil Jae Kim, Seong Jin Cho, and Jae-Wook Jeon. Parallel quick sort algorithms analysis using openmp 3.0 in embedded system. In *2011 11th International Conference on Control, Automation and Systems*, pages 757–761. IEEE, 2011.

## 5 Appendix

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <omp.h>
4 #include <string.h>
5
6
7 /*Reference: Quicksor.pdf, Lab5-lab10, Assignment2-4,
   chatgpt, deepseek,
8 Pacheco, P. S. (2011). An introduction to parallel
   programming. Elsevier/Morgan Kaufmann.
9 Nugroho, E. D., Ashari, I. F., Nashrullah, M., Algifari, M.
   H., & Verdiana, M. (2023). Comparative Analysis of OpenMP
   and MPI Parallel Computing Implementations in Team Sort
   Algorithm. Journal of Applied Informatics and Computing,
   7(2), 141–149 . https://doi.org/10.30871/jaic.v7i2.6409
10 Kil Jae Kim, Seong Jin Cho, & Jae-Wook Jeon. (2011).
   Parallel quick sort algorithms analysis using OpenMP 3.0
   in embedded system. 2011 11th International Conference on
   Control, Automation and Systems, 757–761 .
11 https://www.geeksforgeeks.org/dsa/implementation-of-quick-
   sort-using-mpi-omp-and-posix-thread/
```

```

12  */
13
14  typedef struct {
15      int* data;
16      int size;
17  } ThreadData;
18
19  int compare(const void *a, const void *b) {
20      return (*(int*)a - *(int*)b);
21  }
22
23  void sort_array(int* arr, int size) {
24      qsort(arr, size, sizeof(int), compare);
25  }
26
27  int** global_medians = NULL;
28  int* global_pivots = NULL;
29  int* splitpoints = NULL;
30  int* exchange_sizes = NULL;
31
32  //Here I used double buffer system, I learn it from chatgpt
33  //and lab6(Temporal locality)
34  int** thread_temp_buffers = NULL;
35  int** thread_buffers_A = NULL;
36  int** thread_buffers_B = NULL;
37  // This buffer can check which buffer is active, and buffer
38  // A is 0, buffer B is 1
39  int* buffer_track = NULL;
40  // threadprivate, avoiding race condition
41  static int* findsplit_local_temp_buffer = NULL;
42  static int findsplit_local_temp_buffer_size = 0;
43  #ifdef _OPENMP
44  #pragma omp threadprivate(findsplit_local_temp_buffer,
45  findsplit_local_temp_buffer_size)
46  #endif
47
48  int select_pivot(ThreadData* threadData, int num_threads,
49  int group_size) {
50      int myid = omp_get_thread_num();
51      int locid = myid % group_size;
52      int group = myid / group_size;
53
54      int* data = threadData[myid].data;
55      int size = threadData[myid].size;
56
57      sort_array(data, size);
58
59      int median;
60      if (size == 0){
61          median = 0;

```

```

58     }else if (size == 1) {
59         median = data[0];
60     }else if (size % 2 == 0) {
61         median = (data[size/2 - 1] + data[size/2]) / 2;
62     }else{
63         median = data[size/2];
64     }
65
66     global_medians[group][locid] = median;
67     #pragma omp barrier
68
69     if (locid == 0) {
70         sort_array(global_medians[group], group_size);
71         int mid = group_size / 2;
72         if (group_size % 2 == 0){
73             global_pivots[group]=(global_medians[group][mid
74                                     -1] + global_medians[group][mid]) / 2;
75         }else{
76             global_pivots[group]=global_medians[group][mid];
77         }
78     }
79     #pragma omp barrier
80     return global_pivots[group];
81 }
82
83 int findsplit(int* data, int size, int pivot) {
84     int* temp = findsplit_local_temp_buffer;
85     int left = 0;
86     for (int i = 0; i < size; i++) {
87         if (data[i] <= pivot) {
88             temp[left++] = data[i];
89         }
90     }
91     int split = left;
92     for (int i = 0; i < size; i++) {
93         if (data[i] > pivot) {
94             temp[left++] = data[i];
95         }
96     }
97     memcpy(data, temp, size * sizeof(int));
98     return split;
99 }
100 void merge(int* dest, int* a, int size_a, int* b, int size_b
101 ) {
102     int i = 0, j = 0, k = 0;
103     while (i < size_a && j < size_b)
104         dest[k++] = (a[i] <= b[j]) ? a[i++] : b[j++];
105     while (i < size_a) dest[k++] = a[i++];
106     while (j < size_b) dest[k++] = b[j++];

```



```

106 }
107 /*Here, the data exchange algorithm is inspired MPI
108 1.Each thread has a partner thread for data exchange
109 2.Data is copied into temporary buffers (thread_temp_buffers
    ), mimicking MPI's send/receive buffers.
110 3.shared memory for direct buffer access, avoiding the
    overhead of communication.
111 */
112 void exchange_data(ThreadData* threadData, int group_size) {
113     int myid = omp_get_thread_num();
114     int locid = myid % group_size;
115
116     int* current_data = threadData[myid].data;
117     int current_size = threadData[myid].size;
118     int my_split = splitpoints[myid];
119
120     int *keep_src, *send_src;
121     int keep_size, send_size;
122
123     // For lower group, keeping the data less than pivot,
    and send the data larger than pivot
124     if (locid < group_size / 2) {
125         keep_src = current_data;
126         keep_size = my_split;
127         send_src = current_data + my_split;
128         send_size = current_size - my_split;
129     } else {
130         // For upper group, keeping the data larger than
    pivot, and send the data less than pivot
131         keep_src = current_data + my_split;
132         keep_size = current_size - my_split;
133         send_src = current_data;
134         send_size = my_split;
135     }
136
137     // Copy data to be sent to temporary buffer
138     memcpy(thread_temp_buffers[myid], send_src, send_size *
    sizeof(int));
139     exchange_sizes[myid] = send_size;
140
141     #pragma omp barrier
142
143     // Partner thread for exchange
144     int partner = (locid < group_size / 2) ? myid +
    group_size / 2 : myid - group_size / 2;
145     int recv_size = exchange_sizes[partner];
146     int* recv_data_from_partner = thread_temp_buffers[
    partner]; // Partner's temp buffer
147
148     //target buffer(buffer A or buffer B)

```

```

149     int target_buffer_idx = 1 - buffer_track[myid];
150     int* target_buffer = (target_buffer_idx == 0) ?
        thread_buffers_A[myid] : thread_buffers_B[myid];
151
152
153     merge(target_buffer, keep_src, keep_size,
        recv_data_from_partner, recv_size);
154
155
156     threadData[myid].data = target_buffer;
157     threadData[myid].size = keep_size + recv_size;
158     buffer_track[myid] = target_buffer_idx;
159
160     #pragma omp barrier
161 }
162
163
164 void global_sort(ThreadData* threadData, int
    current_group_size) {
165     if (current_group_size <= 1) {
166         int myid = omp_get_thread_num();
167         sort_array(threadData[myid].data, threadData[myid].
            size);
168         return;
169     }
170
171     int pivot = select_pivot(threadData, omp_get_num_threads
        (), current_group_size);
172
173     int myid = omp_get_thread_num();
174     splitpoints[myid] = findsplit(threadData[myid].data,
        threadData[myid].size, pivot);
175
176     exchange_data(threadData, current_group_size);
177     global_sort(threadData, current_group_size / 2);
178 }
179
180 int main(int argc, char** argv) {
181     if (argc != 3) {
182         printf("Usage: %s <array_size> <num_threads>\n",
            argv[0]);
183         return 1;
184     }
185     int n = atoi(argv[1]);
186     int num_threads = atoi(argv[2]);
187
188     if ((num_threads & (num_threads - 1)) != 0 ||
        num_threads == 0) {
189         printf("Number of threads must be a power of 2 and
            greater than 0.\n");

```

```

190         return 1;
191     }
192
193     omp_set_num_threads(num_threads);
194
195
196     int* data = malloc(n * sizeof(int));
197
198     // Fixed seed for reproducibility, learn from chatgpt
199     srand(42);
200     for (int i = 0; i < n; i++) {
201         data[i] = rand() % 10000;
202     }
203
204
205     ThreadData* threadData = malloc(num_threads * sizeof(
        ThreadData));
206
207     global_pivots = malloc((num_threads / 2) * sizeof(int));
208
209     global_medians = malloc((num_threads / 2) * sizeof(int*)
        );
210
211     for (int i = 0; i < num_threads / 2; ++i) {
212         global_medians[i] = malloc(num_threads * sizeof(int)
            );
213     }
214
215
216     splitpoints = malloc(num_threads * sizeof(int));
217     exchange_sizes = malloc(num_threads * sizeof(int));
218     thread_temp_buffers = malloc(num_threads * sizeof(int*))
        ;
219     thread_buffers_A = malloc(num_threads * sizeof(int*));
220     thread_buffers_B = malloc(num_threads * sizeof(int*));
221     buffer_track = malloc(num_threads * sizeof(int));
222
223
224     double start_time, end_time; // Declare timing variables
        here
225
226     // Each thread initializes its own buffers and calls
        global_sort
227     #pragma omp parallel
228     {
229         int tid = omp_get_thread_num();
230         thread_buffers_A[tid] = malloc(n * sizeof(int));
231         thread_buffers_B[tid] = malloc(n * sizeof(int));
232         thread_temp_buffers[tid] = malloc(n * sizeof(int));
        // Max possible size for temp buffer

```

```

233     findsplit_local_temp_buffer = malloc(n * sizeof(int)
234     );
235     findsplit_local_temp_buffer_size = n;
236     // Start with A as active
237     buffer_track[tid] = 0;
238
239     int chunk_size = n / num_threads;
240     int offset = tid * chunk_size;
241     int current_thread_actual_size = (tid == num_threads
242     - 1) ? (n - offset) : chunk_size;
243     // Copy data to this thread's buffer A
244     memcpy(thread_buffers_A[tid], data + offset,
245     current_thread_actual_size * sizeof(int));
246     threadData[tid].data = thread_buffers_A[tid];
247     threadData[tid].size = current_thread_actual_size;
248
249     #pragma omp barrier
250     // only master thread record runtime
251     #pragma omp master
252     {
253         start_time = omp_get_wtime();
254     }
255
256     global_sort(threadData, num_threads);
257
258     #pragma omp master
259     {
260         end_time = omp_get_wtime();
261     }
262
263     // Each thread frees its own threadprivate buffer,
264     // so we need free it after sort function
265     if (findsplit_local_temp_buffer != NULL) {
266         free(findsplit_local_temp_buffer);
267         findsplit_local_temp_buffer = NULL;
268     }
269
270     // Collect results back into the original data
271     int current_offset = 0;
272     for (int i = 0; i < num_threads; i++) {
273         memcpy(data + current_offset, threadData[i].data,
274         threadData[i].size * sizeof(int));
275         current_offset += threadData[i].size;
276     }
277
278     int is_sorted = 1;
279     for (int i = 1; i < n; i++) {
280         if (data[i-1] > data[i]) {
281             is_sorted = 0;

```

```

278         break;
279     }
280 }
281
282 if (is_sorted) {
283     printf("OK!_Array_is_sorted.\n");
284     printf("Time:_%%.4fs\n", end_time - start_time);
285 } else {
286     printf("Error!_Array_is_not_sorted.\n");
287     return 1;
288 }
289
290
291 for (int i = 0; i < num_threads; i++) {
292     free(thread_buffers_A[i]);
293     free(thread_buffers_B[i]);
294     free(thread_temp_buffers[i]);
295 }
296 free(thread_buffers_A);
297 free(thread_buffers_B);
298 free(thread_temp_buffers);
299
300 for (int i = 0; i < num_threads / 2; ++i) {
301     free(global_medians[i]);
302 }
303 free(global_medians);
304 free(buffer_track);
305 free(splitpoints);
306 free(exchange_sizes);
307 free(global_pivots);
308 free(threadData);
309 free(data);
310
311 return 0;
312 }

```